

# dog\_app

March 2, 2019

## 1 Convolutional Neural Networks

### 1.1 Project: Write an Algorithm for a Dog Identification App

---

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

**Note:** Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

## Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets: \* Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog\_images.

- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

*Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.*

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays `human_files` and `dog_files`.

```
In [28]: import numpy as np
         from glob import glob

         # load filenames for human and dog images
         human_files = np.array(glob("/data/lfw/*/"))
         dog_files = np.array(glob("/data/dog_images/*/"))

         # print number of images in each dataset
         print('There are %d total human images.' % len(human_files))
         print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

### ## Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the `haarcascades` directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [29]: import cv2
         import matplotlib.pyplot as plt
         %matplotlib inline

         # extract pre-trained face detector
         face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

         # load color (BGR) image
         img = cv2.imread(human_files[0])
         # convert BGR image to grayscale
         gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

         # find faces in image
         faces = face_cascade.detectMultiScale(gray)

         # print number of faces detected in the image
         print('Number of faces detected:', len(faces))

         # get bounding box for each detected face
         for (x,y,w,h) in faces:
             # add bounding box to color image
             cv2.rectangle(img, (x,y), (x+w,y+h), (255,0,0), 2)
```

```
# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()
```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

### 1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [30]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
```

```

img = cv2.imread(img_path)
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
faces = face_cascade.detectMultiScale(gray)
return len(faces) > 0

```

### 1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

**Question 1:** Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

**Answer:**

What percentage of the first 100 images in `human_files` have a detected human face? > 98%

What percentage of the first 100 images in `dog_files` have a detected human face? > 17%

In [4]: `from tqdm import tqdm`

```

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

```

```

#-#-# Do NOT modify the code above this line. #-#-#

```

```

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.

```

```

In [41]: #print(human_files_short[0])
         #print(face_detector(human_files_short[0]))
         human_face_found = 0
         for h_imgpath in human_files_short:
             if face_detector(h_imgpath):
                 human_face_found += 1

         print("{}% human faces have been detected correctly from the human_files images.".format(

         dog_face_found = 0
         for d_imgpath in dog_files_short:
             if face_detector(d_imgpath):
                 dog_face_found += 1

         print("{}% human faces have been detected incorrectly in the dog_files images.".format(

```

98.0% human faces have been detected correctly from the `human_files` images.

17.0% human faces have been detected incorrectly in the `dog_files` images.

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [ ]: ### (Optional)
        ### TODO: Test performance of another face detection algorithm.
        ### Feel free to use as many code cells as needed.
```

---

### ## Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

#### 1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```
In [31]: import torch
        import torchvision.models as models

        # check if CUDA is available
        use_cuda = torch.cuda.is_available()

In [32]: # define VGG16 model
        VGG16 = models.vgg16(pretrained=True)

        # move model to GPU if CUDA is available
        if use_cuda:
            VGG16 = VGG16.cuda()
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

```
In [17]: print(VGG16)
```

```
VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
```

```

(5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(6): ReLU(inplace)
(7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(8): ReLU(inplace)
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): ReLU(inplace)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(13): ReLU(inplace)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): ReLU(inplace)
(16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(18): ReLU(inplace)
(19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(25): ReLU(inplace)
(26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(27): ReLU(inplace)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace)
(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace)
  (2): Dropout(p=0.5)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace)
  (5): Dropout(p=0.5)
  (6): Linear(in_features=4096, out_features=1000, bias=True)
)
)

```

#### 1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher\_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```

In [33]: from PIL import Image
import torchvision.transforms as transforms

def VGG16_predict(img_path):
    '''
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    '''

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image

    im = Image.open(img_path)

    #printing the image for a test
    #plt.imshow(im)
    #plt.show()

    normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                     std=[0.229, 0.224, 0.225])

    preprocessing = transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        normalize
    ])

    getImage = transforms.ToPILImage()

    im_preproc = preprocessing(im)

    #printing the image for a test
    #plt.imshow(getImage(im_preproc))
    #plt.show()

    im_preproc.unsqueeze_(0)
    if use_cuda:

```

```

        im_preproc = im_preproc.cuda()

    output = VGG16(im_preproc)

    if use_cuda:
        output = output.cpu()

    predicted_Index = output.data.numpy().argmax()

    #max index
    #print(predicted_Index)

    return predicted_Index # predicted class index

#VGG16_predict(human_files[1])

```

### 1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the dog\_detector function below, which returns True if a dog is detected in an image (and False if not).

```

In [8]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.
    pred_Index = VGG16_predict(img_path)
    #print(pred_Index)
    if pred_Index >=151 and pred_Index <=268:

        return True
    else:
        return False

#return None # true/false

print("This is a Human. Dog ditector:{}".format(dog_detector(human_files[1])))
print("This is a dog. Dog ditector:{}".format(dog_detector(dog_files[1])))
print("This is a Human. Wrongly classified. Dog ditector:{}".format(dog_detector('/data

This is a Human. Dog ditector:False
This is a dog. Dog ditector:True
This is a Human. Wrongly classified. Dog ditector:False

```



### 1.1.6 (IMPLEMENTATION) Assess the Dog Detector

**Question 2:** Use the code cell below to test the performance of your dog\_detector function.

- What percentage of the images in human\_files\_short have a detected dog?
- What percentage of the images in dog\_files\_short have a detected dog?

**Answer:**

What percentage of the images in human\_files\_short have a detected dog? > 1%

What percentage of the images in dog\_files\_short have a detected dog? > 100%

```
In [42]: ### TODO: Test the performance of the dog_detector function  
        ### on the images in human_files_short and dog_files_short.
```

```
human_face_found = 0  
for h_imgpath in human_files_short:  
    if dog_detector(h_imgpath):  
        human_face_found += 1  
  
print("{}% dog faces have been detected from the human_files images.".format(human_face_found))  
  
dog_face_found = 0  
for d_imgpath in dog_files_short:  
    if dog_detector(d_imgpath):  
        dog_face_found += 1  
  
print("{}% dog faces have been detected in the dog_files images.".format(dog_face_found))
```


1.0% dog faces have been detected from the human\_files images.

100.0% dog faces have been detected in the dog\_files images.

```
In [44]: #printing the error detection image for a test
```

```
for h_imgpath in human_files_short:  
    if dog_detector(h_imgpath):  
        print(h_imgpath)  
        im = Image.open(h_imgpath)  
        plt.imshow(im)  
        plt.show()  
        break
```

/data/lfw/Perri\_Shaw/Perri\_Shaw\_0001.jpg



We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [ ]: ### (Optional)
        ### TODO: Report the performance of another pre-trained network.
        ### Feel free to use as many code cells as needed.
```

---

### ## Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

---

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

---

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

---

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

---

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

### 1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [34]: import os
         from torchvision import datasets

         #otherwise throwing error during conversions (training)
         from PIL import ImageFile
         ImageFile.LOAD_TRUNCATED_IMAGES = True

         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes

         batch_size = 20
         test_batch_size = 15

         normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                         std=[0.229, 0.224, 0.225])

         train_transforms_process = transforms.Compose([
             transforms.RandomResizedCrop(224),
             transforms.RandomHorizontalFlip(),
             transforms.RandomRotation(10),
             transforms.ToTensor(),
             normalize
         ])
```

```

transforms_process = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.ToTensor(),
    normalize
])

train_data = datasets.ImageFolder('/data/dog_images/train', transform=train_transforms_)
valid_data = datasets.ImageFolder('/data/dog_images/valid', transform=transforms_process)
test_data = datasets.ImageFolder('/data/dog_images/test', transform=transforms_process)

train_data_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size, shuffle=True)
valid_data_loader = torch.utils.data.DataLoader(valid_data, batch_size=test_batch_size)
test_data_loader = torch.utils.data.DataLoader(test_data, batch_size=test_batch_size)

loaders_scratch ={'train': train_data_loader,
                  'valid': valid_data_loader,
                  'test': test_data_loader}

In [35]: print(" train data: {}".format(len(train_data)))
        print(" valid data: {}".format(len(valid_data)))
        print(" test data: {}".format(len(test_data)))

train data: 6680
valid data: 835
test data: 836

```

**Question 3:** Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

**Answer:** >For the training dataset, I wanted to create  $224 * 224 * 3$  tensors. So I randomly cropped the image in  $224 * 224$  dimension. This was big enough image size to work on the current problem.

Yes, the image was cropped in the above size, rotated 10 degree randomly and horizontally flipped randomly. This was done to avoid overfitting of data and also to make the model more efficient for cases when it gets slight distorted images.

### 1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [36]: import torch.nn as nn
        import torch.nn.functional as F

        # define the CNN architecture

```

```

class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN

        self.conv1 = nn.Conv2d(3, 16, 3, padding=1) # 224 * 224 * 3 > 112 * 112 * 16
        self.conv2 = nn.Conv2d(16, 32, 3, padding=1) # 112 * 112 * 16 > 56 * 56 * 64
        self.conv3 = nn.Conv2d(32, 64, 3, padding=1) # 56 * 56 * 32 > 28 * 28 * 64
        self.conv4 = nn.Conv2d(64, 128, 3, padding=1) # 28 * 28 * 64 > 14 * 14 * 128

        self.pool = nn.MaxPool2d(2,2)

        self.fc1 = nn.Linear(128*14*14, 500)
        self.fc2 = nn.Linear(500,133) #output size 133

        self.dropout = nn.Dropout(0.25)

    def forward(self, x):
        ## Define forward behavior
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))
        x = self.pool(F.relu(self.conv4(x)))

        x = x.view(-1, 128*14*14)

        x= self.dropout(x)
        x = F.relu(self.fc1(x))

        x= self.dropout(x)
        x = self.fc2(x)

        return x

##-## You so NOT have to modify the code below this line. ##-##

# instantiate the CNN
model_scratch = Net()
print(model_scratch)

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

```

```

Net(
  (conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

(conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(fc1): Linear(in_features=25088, out_features=500, bias=True)
(fc2): Linear(in_features=500, out_features=133, bias=True)
(dropout): Dropout(p=0.25)
)

```

**Question 4:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

**Answer:**

After input image processing, I decide the input layer to be  $224 * 224 * 3$  tensor. So in the model, the first convolution layer takes  $224 * 224 * 3$  as input and outputs  $112 * 112 * 16$ . It reduces the dimension by half and increase the depth by 2 times. Likewise all convolution layer gradually increasing the depth and reduces the dimension, this helps the model to extract more features from the input image.

I started with only 3 convolution layers but it did not produce desired accuracy. So I increase the convolution layer and with 4 layers I was able to reach the target.

After convolution layer, MaxPool2d is used to down-sample so that eventually we can flatten the layers. I have used 2 linear layers to flatten the features and narrowed it down. The network is trying to categorise the images into 133 classes so the output layer is one of 133. I have also used dropout (with probability of .25) to avoid overfitting and used Relu in the linear layer as it has been very efficient.

### 1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```

In [37]: import torch.optim as optim

        ### TODO: select loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### TODO: select optimizer
        optimizer_scratch = optim.Adam(model_scratch.parameters())

```

### 1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```

In [38]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
        """returns trained model"""
        # initialize tracker for minimum validation loss
        valid_loss_min = np.Inf

```

```

for epoch in range(1, n_epochs+1):
    # initialize variables to monitor training and validation loss
    train_loss = 0.0
    valid_loss = 0.0

    #####
    # train the model #
    #####
    model.train()
    for batch_idx, (data, target) in enumerate(loaders['train']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
            ## find the loss and update the model parameters accordingly
            ## record the average training loss, using something like
            ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
            optimizer.zero_grad()

        output = model(data)

        loss = criterion(output, target)
        loss.backward()
        optimizer.step()

        train_loss = train_loss + ((1/(batch_idx+1))* (loss.data-train_loss))
    #####
    # validate the model #
    #####
    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['valid']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
            ## update the average validation loss
            output = model(data)
            loss = criterion(output, target)

        valid_loss = valid_loss + ((1/(batch_idx+1))* (loss.data-valid_loss))

    train_loss = train_loss/len(train_data_loader.dataset)
    valid_loss = valid_loss/len(valid_data_loader.dataset)

    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch,
        train_loss,

```

```

        valid_loss
    ))

    ## TODO: save the model if validation loss has decreased

    if valid_loss <= valid_loss_min:
        print("Saving Model. Valid_loss {} is less than valid_loss_min {}".format(
            valid_loss, valid_loss_min))
        torch.save(model.state_dict(), save_path)
        valid_loss_min = valid_loss

    # return trained model
    return model

In [39]: # train the model
#100
model_scratch = train(30, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'model_scratch.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))

Epoch: 1      Training Loss: 0.000727      Validation Loss: 0.005694
Saving Model. Valid_loss 0.005694288294762373 is less than valid_loss_min inf
Epoch: 2      Training Loss: 0.000702      Validation Loss: 0.005456
Saving Model. Valid_loss 0.005455581936985254 is less than valid_loss_min 0.005694288294762373
Epoch: 3      Training Loss: 0.000680      Validation Loss: 0.005286
Saving Model. Valid_loss 0.005286138970404863 is less than valid_loss_min 0.005455581936985254
Epoch: 4      Training Loss: 0.000663      Validation Loss: 0.005142
Saving Model. Valid_loss 0.005141913890838623 is less than valid_loss_min 0.005286138970404863
Epoch: 5      Training Loss: 0.000654      Validation Loss: 0.005155
Epoch: 6      Training Loss: 0.000642      Validation Loss: 0.005033
Saving Model. Valid_loss 0.005032563582062721 is less than valid_loss_min 0.005141913890838623
Epoch: 7      Training Loss: 0.000632      Validation Loss: 0.004924
Saving Model. Valid_loss 0.004923985805362463 is less than valid_loss_min 0.005032563582062721
Epoch: 8      Training Loss: 0.000623      Validation Loss: 0.004870
Saving Model. Valid_loss 0.004870202392339706 is less than valid_loss_min 0.004923985805362463
Epoch: 9      Training Loss: 0.000618      Validation Loss: 0.004891
Epoch: 10     Training Loss: 0.000612      Validation Loss: 0.004781
Saving Model. Valid_loss 0.004781119525432587 is less than valid_loss_min 0.004870202392339706
Epoch: 11     Training Loss: 0.000607      Validation Loss: 0.004767
Saving Model. Valid_loss 0.004766898695379496 is less than valid_loss_min 0.004781119525432587
Epoch: 12     Training Loss: 0.000600      Validation Loss: 0.004770
Epoch: 13     Training Loss: 0.000592      Validation Loss: 0.004680
Saving Model. Valid_loss 0.0046796840615570545 is less than valid_loss_min 0.004766898695379496
Epoch: 14     Training Loss: 0.000586      Validation Loss: 0.004750
Epoch: 15     Training Loss: 0.000581      Validation Loss: 0.004612
Saving Model. Valid_loss 0.0046124448999762535 is less than valid_loss_min 0.0046796840615570545

```



```

Epoch: 16          Training Loss: 0.000577          Validation Loss: 0.004595
Saving Model. Valid_loss 0.004594812635332346 is less than valid_loss_min 0.0046124448999762535
Epoch: 17          Training Loss: 0.000573          Validation Loss: 0.004756
Epoch: 18          Training Loss: 0.000563          Validation Loss: 0.004541
Saving Model. Valid_loss 0.004540877416729927 is less than valid_loss_min 0.004594812635332346
Epoch: 19          Training Loss: 0.000560          Validation Loss: 0.004505
Saving Model. Valid_loss 0.004504651762545109 is less than valid_loss_min 0.004540877416729927
Epoch: 20          Training Loss: 0.000553          Validation Loss: 0.004505
Epoch: 21          Training Loss: 0.000551          Validation Loss: 0.004459
Saving Model. Valid_loss 0.004459007643163204 is less than valid_loss_min 0.004504651762545109
Epoch: 22          Training Loss: 0.000549          Validation Loss: 0.004509
Epoch: 23          Training Loss: 0.000549          Validation Loss: 0.004463
Epoch: 24          Training Loss: 0.000541          Validation Loss: 0.004493
Epoch: 25          Training Loss: 0.000541          Validation Loss: 0.004432
Saving Model. Valid_loss 0.004432493820786476 is less than valid_loss_min 0.004459007643163204
Epoch: 26          Training Loss: 0.000535          Validation Loss: 0.004479
Epoch: 27          Training Loss: 0.000534          Validation Loss: 0.004459
Epoch: 28          Training Loss: 0.000527          Validation Loss: 0.004427
Saving Model. Valid_loss 0.004427027422934771 is less than valid_loss_min 0.004432493820786476
Epoch: 29          Training Loss: 0.000526          Validation Loss: 0.004352
Saving Model. Valid_loss 0.004351572133600712 is less than valid_loss_min 0.004427027422934771
Epoch: 30          Training Loss: 0.000527          Validation Loss: 0.004265
Saving Model. Valid_loss 0.00426473980769515 is less than valid_loss_min 0.004351572133600712

```

### 1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```

In [10]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))

```

```

        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

```

```

In [14]: # call test function
         test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 3.640179

Test Accuracy: 18% (153/836)

---

#### ## Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

##### 1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```

In [11]: ## TODO: Specify data loaders
         import os
         from torchvision import datasets
         import torchvision.transforms as transforms

         #otherwise throwing error during conversions (training)
         from PIL import ImageFile
         ImageFile.LOAD_TRUNCATED_IMAGES = True

         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes

         batch_size = 20
         test_batch_size = 15

```

```

normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                std=[0.229, 0.224, 0.225])

train_transforms_process = transforms.Compose([
    transforms.RandomResizedCrop(224),
    transforms.RandomHorizontalFlip(),
    #transforms.RandomVerticalFlip(),
    transforms.RandomRotation(10),
    transforms.ToTensor(),
    normalize
])

transforms_process = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.ToTensor(),
    normalize
])

train_data = datasets.ImageFolder('/data/dog_images/train', transform=train_transforms_process)
valid_data = datasets.ImageFolder('/data/dog_images/valid', transform=transforms_process)
test_data = datasets.ImageFolder('/data/dog_images/test', transform=transforms_process)

train_data_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size, shuffle=True)
valid_data_loader = torch.utils.data.DataLoader(valid_data, batch_size=test_batch_size)
test_data_loader = torch.utils.data.DataLoader(test_data, batch_size=test_batch_size)

loaders_transfer = {'train': train_data_loader,
                    'valid': valid_data_loader,
                    'test': test_data_loader}

```

### 1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

In [12]: import torchvision.models as models
import torch.nn as nn

## TODO: Specify model architecture
model_transfer = models.resnet18(pretrained=True)

for param in model_transfer.parameters():
    param.requires_grad = False

print(model_transfer)
model_transfer.fc = nn.Linear(512, 133, bias=True)

```

```

    #for param in model_transfer.fc.parameters():
    #    param.required_grad = True

```

```

print(model_transfer)

```

```

if use_cuda:
    model_transfer = model_transfer.cuda()

```

Downloading: "https://download.pytorch.org/models/resnet18-5c106cde.pth" to /root/.torch/models/100%|| 46827520/46827520 [00:01<00:00, 25211542.06it/s]

```

ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer2): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

        (relu): ReLU(inplace)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(avgpool): AvgPool2d(kernel_size=7, stride=1, padding=0)
(fc): Linear(in_features=512, out_features=1000, bias=True)
)

```

```

ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer2): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer3): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
)

```

```

        (downsample): Sequential(
          (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
      )
    (1): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(avgpool): AvgPool2d(kernel_size=7, stride=1, padding=0)
(fc): Linear(in_features=512, out_features=133, bias=True)
)

```

**Question 5:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

**Answer:** Any pretrained torchvision models probably would have been good enough for the given problem. I have chosen ResNet because, based on some online search, I found ResNet has better performance than some of the other models. It even won the ImageNet prize in 2015.

I am using the pretrained model as is, apart from the last linear layer. The model had 1000 output features but in this problem, we need only 133 output features. So I replaced the last linear layer with a new one. We also had to set the `requires_grad` to false so that the model does not back propagate and start update the weights of all other layers for the new training set. Our intention is to train only the last layer of the model.

### 1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [13]: import torch.optim as optim

criterion_transfer = nn.CrossEntropyLoss()

optimizer_transfer = optim.Adam(model_transfer.fc.parameters())
```

### 1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [17]: # train the model
        #model_transfer = # train(n_epochs, loaders_transfer, model_transfer, optimizer_transfer,
        n_epochs = 10

        model_transfer = train(n_epochs, loaders_transfer, model_transfer, optimizer_transfer,

Epoch: 1          Training Loss: 0.000193          Validation Loss: 0.000946
Saving Model. Valid_loss 0.0009458689601160586 is less than valid_loss_min inf
Epoch: 2          Training Loss: 0.000177          Validation Loss: 0.000906
Saving Model. Valid_loss 0.0009062682511284947 is less than valid_loss_min 0.0009458689601160586
Epoch: 3          Training Loss: 0.000170          Validation Loss: 0.000908
Epoch: 4          Training Loss: 0.000164          Validation Loss: 0.000839
Saving Model. Valid_loss 0.0008389318827539682 is less than valid_loss_min 0.0009062682511284947
Epoch: 5          Training Loss: 0.000166          Validation Loss: 0.000859
Epoch: 6          Training Loss: 0.000161          Validation Loss: 0.000848
Epoch: 7          Training Loss: 0.000161          Validation Loss: 0.000916
Epoch: 8          Training Loss: 0.000161          Validation Loss: 0.000877
Epoch: 9          Training Loss: 0.000154          Validation Loss: 0.000897
Epoch: 10         Training Loss: 0.000153          Validation Loss: 0.000854
```

```
In [14]: # load the model that got the best validation accuracy (uncomment the line below)
        model_transfer.load_state_dict(torch.load('model_transfer.pt'))
```

### 1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [18]: test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)

Test Loss: 0.740331
```

Test Accuracy: 78% (653/836)



### 1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [23]: ### TODO: Write a function that takes a path to an image as input
### and returns the dog breed that is predicted by the model.
data_transfer = {'train': train_data,
                 'valid': valid_data,
                 'test': test_data}

# list of class names by index, i.e. a name can be accessed like class_names[0]
class_names = [item[4:].replace("_", " ") for item in data_transfer['train'].classes]

def predict_breed_transfer(img_path):
    # load the image and return the predicted breed
    im = Image.open(img_path)

    #printing the image for a test
    #plt.imshow(im)
    #plt.show()

    normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                     std=[0.229, 0.224, 0.225])

    preprocessing = transforms.Compose([
        transforms.Resize((224,224)),
        transforms.ToTensor(),
        normalize
    ])

    #getImage = transforms.ToPILImage()

    im_preproc = preprocessing(im)
    im_preproc.unsqueeze_(0)

    if use_cuda:
        im_preproc = im_preproc.cuda()

    output = model_transfer(im_preproc)

    #if use_cuda:
    #    output = output.cpu()

    #print(output)
    #print(torch.max(output.data,1))
```



Sample Human Output

```
_, predicted_tensor = torch.max(output,1)

if use_cuda:
    predicted_Index = np.squeeze(predicted_tensor.cpu().numpy())
else:
    predicted_Index = np.squeeze(predicted_tensor.numpy())

print("Predicted Index {}".format(predicted_Index))
dog_breed = class_names[predicted_Index]
return dog_breed
```

```
In [24]: #print(dog_files_short[0])
         #print(human_files_short[0])

         predict_breed_transfer(dog_files_short[0])
```

Predicted Index 40

Out[24]: 'Bullmastiff'

---

### ## Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

### 1.1.18 (IMPLEMENTATION) Write your Algorithm

```
In [25]: import matplotlib.pyplot as plt
        from PIL import Image

        ### TODO: Write your algorithm.
        ### Feel free to use as many code cells as needed.

        def run_app(img_path):
            ## handle cases for a human face, dog, and neither

            im = Image.open(img_path)
            #printing the image for a test
            plt.imshow(im)
            plt.show()

            if face_detector(img_path) == True:
                prediction = predict_breed_transfer(img_path)
                print("Human detected. Looks like: {}".format(prediction))
            elif dog_detector(img_path) == True:
                prediction = predict_breed_transfer(img_path)
                print("Dog detected. Looks like: {}".format(prediction))
            else:
                print("Could not detect any dog or human in the picture.")
```

---

#### ## Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

### 1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

**Question 6:** Is the output better than you expected :) ? Or worse :( ? Provide at least three possible points of improvement for your algorithm.

#### Answer:


I was fascinated that with such little training, the model can detect if there is no human/dog in the image and the breed of dogs. With only 45 Mb memory it can detect dog breeds more accurately than me!

There are several things can be done to improve the performance. Here are three: \* Use more varieties of training dataset to train the model \* Use some other models and compare the result. Pick the model that performs the best \* Increase the number of epoch and training time when we can keep reducing the validation loss.

```
In [26]: ## TODO: Execute your algorithm from Step 6 on
        ## at least 6 images on your computer.
```


```
## Feel free to use as many code cells as needed.

## suggested code, below
for file in np.hstack((human_files[:6], dog_files[:6])):
    run_app(file)
```




Predicted Index 64

Human detected. Looks like: Entlebucher mountain dog




Predicted Index 55  
Human detected. Looks like: Dachshund



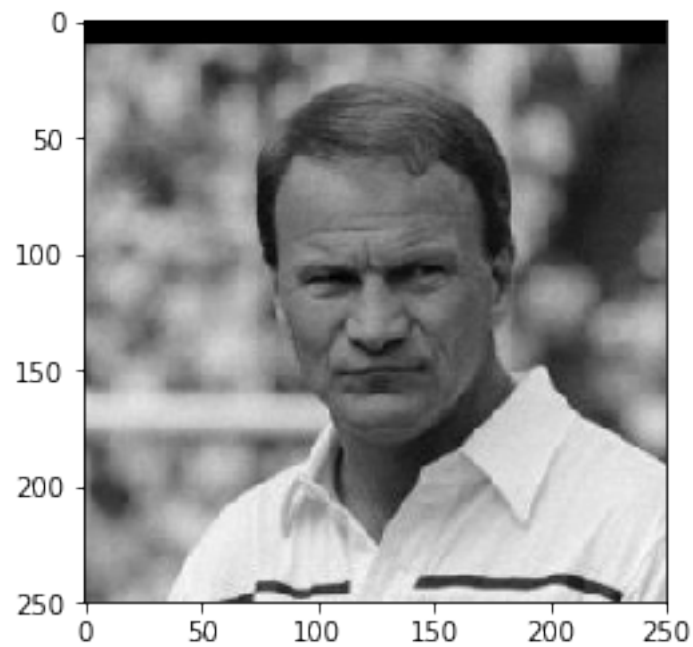
Predicted Index 8

Human detected. Looks like: American water spaniel




Predicted Index 96

Human detected. Looks like: Lakeland terrier




Predicted Index 123

Human detected. Looks like: Poodle




Predicted Index 13

Human detected. Looks like: Basenji




Predicted Index 40  
Dog detected. Looks like: Bullmastiff






Predicted Index 9

Dog detected. Looks like: Anatolian shepherd dog




Predicted Index 40

Dog detected. Looks like: Bullmastiff



Predicted Index 102

Dog detected. Looks like: Mastiff



Predicted Index 40

Dog detected. Looks like: Bullmastiff



Predicted Index 102  
Dog detected. Looks like: Mastiff

```
In [27]: import numpy as np
         from glob import glob
         import matplotlib.pyplot as plt
         from PIL import Image


         # load filenames for human and dog images
         test_files = np.array(glob('test_images/*'))

         # print number of images in each dataset
         print('There are %d total test images.' % len(test_files))
         for img_path in test_files:
             run_app(img_path)
             #im = Image.open(img_path)
             #printing the image for a test
             #plt.imshow(im)
             #plt.show()
```

There are 7 total test images.




Predicted Index 19  
Dog detected. Looks like: Belgian malinois




Predicted Index 55

Human detected. Looks like: Dachshund




Could not detect any dog or human in the picture.




Predicted Index 59

Human detected. Looks like: Dogue de bordeaux



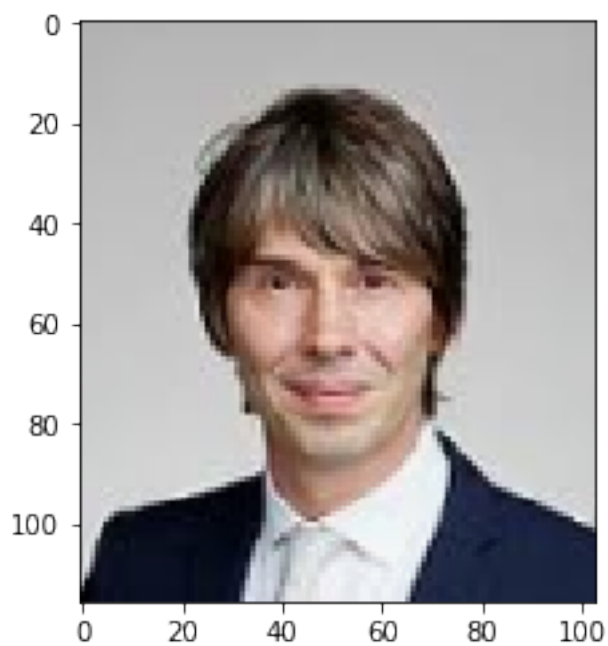
Predicted Index 123

Human detected. Looks like: Poodle



Predicted Index 59

Dog detected. Looks like: Dogue de bordeaux



```
Predicted Index 51  
Human detected. Looks like: Clumber spaniel
```

```
In [ ]:
```